

Gibbs sampling by completion

Oswaldo Gressani

October 12, 2023

Consider the example of [Robert \(1996\)](#), where the Gibbs sampler is illustrated with a completion technique. Completion of a density f by a density g means that the latter density is constructed such that f appears as a marginal density of g . Assume we want to sample from the following density:

$$p(\theta|\theta_0) \propto \frac{\exp(-\theta^2/2)}{(1 + (\theta - \theta_0)^2)^\nu}, \quad \theta \in \mathbb{R}, \nu > 0, \quad (1)$$

where θ_0 is a known real number. To find the density g that completes $p(\theta|\theta_0)$, note that $p(\theta|\theta_0)$ can be written by means of the following integral:

$$\begin{aligned} p(\theta|\theta_0) &\propto \int_0^{+\infty} \exp(-\theta^2/2) \exp(-(1 + (\theta - \theta_0)^2)\eta/2) \eta^{\nu-1} d\eta \\ &\propto \exp(-\theta^2/2) \int_0^{+\infty} \eta^{\nu-1} \exp\left(-\eta \left(\frac{1 + (\theta - \theta_0)^2}{2}\right)\right) d\eta \\ &\propto \exp(-\theta^2/2) \Gamma(\nu) \left(\frac{1 + (\theta - \theta_0)^2}{2}\right)^{-\nu} \\ &\propto \frac{\exp(-\theta^2/2)}{(1 + (\theta - \theta_0)^2)^\nu}. \end{aligned}$$

In other words, the integrand corresponding to the joint density $g(\theta, \eta) \propto \exp(-\theta^2/2) \exp(-(1 + (\theta - \theta_0)^2)\eta/2) \eta^{\nu-1}$ with $\eta > 0$ completes $p(\theta|\theta_0)$. From here, we can easily obtain the conditional of η given θ :

$$\begin{aligned} g_1(\eta|\theta) &= \frac{(0.5(1 + (\theta - \theta_0)^2))^\nu}{\Gamma(\nu)} \eta^{\nu-1} \exp(-(1 + (\theta - \theta_0)^2)\eta/2) \\ \Rightarrow (\eta|\theta) &\sim \mathcal{G}(\nu, 0.5(1 + (\theta - \theta_0)^2)), \end{aligned} \quad (2)$$

where $\mathcal{G}(a, b)$ is a Gamma distribution with shape parameter $a > 0$ and rate parameter $b > 0$, with mean a/b and variance a/b^2 . We can also show that the conditional distribution of θ given η is given by:

$$g_2(\theta|\eta) = \frac{\sqrt{1+\eta}}{\sqrt{2\pi}} \exp\left(-\frac{(1+\eta)}{2} \left(\theta - \frac{\theta_0\eta}{1+\eta}\right)^2\right)$$

$$(\theta|\eta) \sim \mathcal{N}\left(\frac{\theta_0\eta}{1+\eta}, \frac{1}{1+\eta}\right). \quad (3)$$

From (2) and (3), one can easily build a Gibbs sampler to sample from $p(\theta|\theta_0)$.

```
#-----
#
# Gibbs sampling by completion
# Copyright, Oswaldo Gressani. All rights reserved.
#-----

theta0 <- 2
nu <- 3

#--- Target distribution to sample from
ptarget <- function(theta){
  val <- exp(-0.5 * theta^2)/((1+(theta-theta0)^2)^nu)
  return(val)
}

#--- Plot target
x <- seq(-1,4, length = 1000)
dx <- x[2]-x[1]
target_val <- sapply(x, ptarget)
cnorm <- 1/(sum(target_val * dx))
target_val <- cnorm * target_val
#sum(target_val * dx)
plot(x, target_val, type="l", xlab=expression(theta),
ylab="Posterior", lwd = 2)

#-- Gibbs sampler
theta_sample <- c()
eta_sample <- c()
theta_init <- 1
M <- 10000

for(m in 1:M){
  eta_sample[m] <- rgamma(n = 1, shape = nu,
rate = 0.5 * (1+(theta_init-theta0)^2))
  theta_sample[m] <- rnorm(n = 1,
mean = theta0 * (eta_sample[m]/(1+eta_sample[m])),
sd = sqrt(1/(1+eta_sample[m])))
```

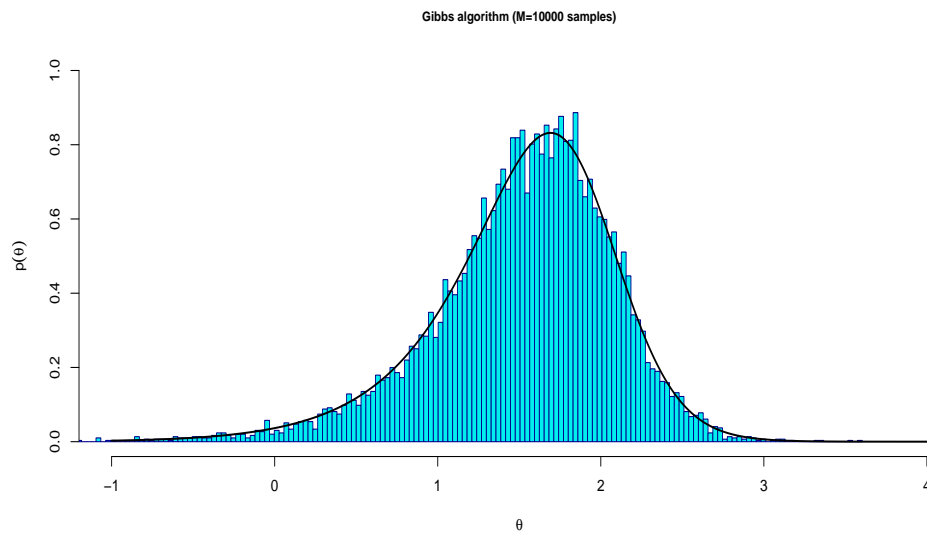
```

theta_init <- theta_sample[m]
}

hist(theta_sample,
breaks=seq(min(theta_sample),max(theta_sample),length.out = 200),
freq = FALSE,
xlim = c(-1,4),
ylim=c(0,1),
ylab=expression(p(theta)),
xlab=expression(theta),
col="cyan2",
border = "darkblue",
main="Gibbs algorithm (M=10000 samples)",
cex.main=0.8)
lines(x, target_val, type="l", lwd = 2)

summary(theta_sample)

```



References

Robert, C. Méthodes de Monte Carlo par chaînes de Markov. Economica, 1996.